

Spark Streaming Assignment: Real-Time Advertisement Data Aggregation

Objective: Process real-time advertisement data using Spark Streaming to gain business insights and store the aggregated data into Cassandra.

Background:

You have been provided with a Kafka topic named ads_data that contains advertisement data in the following format:

```
{  
  "ad_id": "12345",  
  "timestamp": "2023-08-23T12:01:05Z",  
  "clicks": 5,  
  "views": 10,  
  "cost": 50.75  
}
```

The goal is to process this real-time data, compute business insights using window-based aggregation, and write the aggregated results into a Cassandra table. The aggregation key is ad_id, and aggregated values should update previous values in the Cassandra table.

Tasks:

- **Kafka setup and Mock data producer:**
 - Set up Confluent Kafka on cloud or local
 - Create topic named as **ads_data**
 - Write a python script which will use above mentioned data format and keep on publishing random mock data in avro serialized form into Kafka topic
- **Reading Data from Kafka:**
 - Set up a Spark Streaming application.
 - Use the Kafka connector to read data from the ads_data topic.
 - Parse & deserialize the incoming data into the appropriate structure.
- **Windowing Based Aggregation:**

- Perform a window-based aggregation over a window duration (e.g., 1 minute) and sliding interval (e.g., 30 seconds).
- Aggregate the following:
 - Total clicks per ad_id.
 - Total views per ad_id.
 - Average cost per view for each ad_id.
- **Write Aggregated Data to Cassandra:**
 - For each ad_id, check if an entry already exists in the Cassandra table.
 - If an entry exists, update the values:
 - Add new clicks/views to the existing counts.
 - Update the average cost per view.
 - If an entry doesn't exist, create a new row with the aggregated values.

Submission:

Submit your Spark Streaming application code, along with a brief report detailing the results and any challenges faced during the assignment.